SQL Interview Question #1

Let's say you have two SQL tables: `authors` and `books`.

The `authors` dataset has 1M+ rows; here's the first six rows:

author_name	book_name
author_1	book_1
author_1	book_2
author_2	book_3
author_2	book_4
author_2	book_5
author_3	book_6
...	...

The `books` dataset also has 1M+ rows and here's the first six:

book_name	sold_copies
book_1	1000
book_2	1500
book_3	34000
book_4	29000
book_5	40000
book_6	4400
...	...

Create an SQL query that shows the TOP 3 authors who sold the most books in total!

Answer:

```
SELECT author_name
FROM authors, books
WHERE authors.book_name = books.book_name
GROUP BY author_name
ORDER BY sum(sold_copies) DESC
LIMIT 3
```

SQL Interview Question #2

You work for a startup that makes an online presentation software. You have an event log that records every time a user inserted an image into a presentation. (One user can insert many images.) The `event_log` SQL table looks like this:

user_id	event_date_time
7494212	1535308430
7494212	1535308433
1475185	1535308444
6946725	1535308475
6946725	1535308476
6946725	1535308477
...	...

…and it has over one billion rows.

Note: The `event_date_time` *column's format is an epoch timestamp*

Write an SQL query to find out how many users inserted more than 1000 but less than 2000 images in their presentations!

Answer:

```
SELECT count(user_id)
FROM event_log
GROUP BY user_id
HAVING count(event_data_time) >1000 and
count(event_data_time) < 2000
```

SQL Interview Question #3

You have two SQL tables! The first one is called `employees` and it contains the employee names, the unique employee ids and the department names of a company. Sample:

department_name	employee_id	employee_name
Sales	123	John Doe
Sales	211	Jane Smith
HR	556	Billy Bob
Sales	711	Robert Hayek
Marketing	235	Edward Jorgson
Marketing	236	Christine Packard
...	...	...

The second one is named `salaries`. It holds the same employee names and the same employee ids – and the salaries for each employee. Sample:

salary	employee_id	employee_name
500	123	John Doe
600	211	Jane Smith
1000	556	Billy Bob
400	711	Robert Hayek
1200	235	Edward Jorgson
200	236	Christine Packard
...	...	...

The company has 546 employees, so both tables have 546 rows.

Print every department where the average salary per employee is lower than $500!

Answer:

```
SELECT department_name
FROM employees, salaries
WHERE employees.employee_id = salaries.employee_id
and employees.employee_name = salaries. employee_name
GROUP BY department_name
HAVING avg(salary) < 500
```